


The Perceptron and Adaline is one of the oldest and simplest learning classifiers out there

What Adaline and the Perceptron have in common

- they are classifiers for binary classification
- both have a linear decision boundary
- both can learn iteratively, sample by sample (the Perceptron naturally, and Adaline via stochastic gradient descent)
- both use a threshold function




Example of a linear decision boundary for binary classification.

The first step in the two algorithms is to compute the so-called net input z as the linear combination of our feature variables x and the model weights w .

$$\begin{aligned} \mathbf{z} &= w_1x_1 + \dots + w_mx_m = \sum_{j=1}^m x_jw_j \\ &= \mathbf{w}^T \mathbf{x} \end{aligned}$$

Then, in the Perceptron and Adaline, we define a threshold function to make a prediction. I.e., if z is greater than a threshold θ usually 0, we predict class 1, and 0 otherwise:

$$g(\mathbf{z}) = \begin{cases} 1 & \text{if } \mathbf{z} \geq \theta \\ -1 & \text{otherwise.} \end{cases}$$



The difference between the 2 algorithms is that the Perceptron uses the result of the threshold function $g(z)$ to learn the model weights while Adaline uses the net input z (where the activation function is linear) to learn the model coefficients as illustrated below

Perceptron Learning Algorithm

The perceptron learning algorithm is done sample by sample

1. Initialize weight parameters to initial values (usually small random values, rather than zero).
2. Calculate z
3. Predict of the binary output from the given features for each training sample input using the threshold function $\hat{y} = g(z)$.
4. Update of the weights through the Perceptron Learning Rule.

$$w_j := w_j + \alpha(\text{actual output} - \text{predicted output})x_j^{(i)}$$

$$w_j := w_j + \alpha(y^{(i)} - \hat{y}^{(i)})x_j^{(i)}$$

Adaline Learning Process

The Adaline learning process is done sample by sample or on all dataset. It uses gradient descent optimization function

1. Initialize weight parameters to initial values (usually small random values, rather than zero).
2. calculate z
3. apply linear activation function where $\Phi(z) = z$
4. Update of the weights using gradient descent optimization method

$$\Delta w_j = -\alpha \frac{\partial J}{\partial w_j}$$

$$\Delta w_j = \alpha \sum_i (y^{(i)} - \phi(z^{(i)}))x_j^{(i)}$$

Where ∇ is used to denote the gradient, and α is the learning rate.

In []:

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn import datasets
from sklearn.preprocessing import StandardScaler
```

In [2]: **class** Perceptron:

```

def __init__(self, learning_rate=0.0001 , epochs= 1000):
    self.weights = None
    self.bias = None
    self.learning_rate = learning_rate
    self.epochs = epochs

# heaviside activation function
def activation(self, z):
    return np.heaviside(z, 0) # heaviside(z) heaviside -> activation

def fit(self, x, y):
    n_features = x.shape[1]

    # Initializing weights and bias
    self.weights = np.zeros((n_features))
    self.bias = 0

    # Iterating until the number of epochs
    for epoch in range(self.epochs):
        # Traversing through the entire training set
        for i in range(len(x)):
            y_pred = 0
            for j in range(2):
                y_pred += x.iloc[i, j] * self.weights[j]

            y_pred += self.bias

            if y_pred >= 0:
                y_pred = 1
            else:
                y_pred = 0

            for j in range(2):
                self.weights[j] += self.learning_rate * (y.iloc[i] - y_pred)
            self.bias += self.learning_rate * (y.iloc[i] - y_pred)
        return self.weights, self.bias

def predict(self, X):
    z = np.dot(X, self.weights) + self.bias
    return self.activation(z)

```

```
In [3]: class Adaline:
    def __init__(self, learning_rate=0.01, num_epochs=1000):
        self.learning_rate = learning_rate
        self.num_epochs = num_epochs
        self.weights = None
        self.bias = None

    def fit(self, X, y):
        num_samples, num_features = X.shape
        self.weights = np.zeros(num_features)
        self.bias = 0
        error = np.zeros(self.num_epochs)

        for epoch in range(self.num_epochs):
            net_input = np.dot(X, self.weights) + self.bias
            output = self.activation(net_input)
            error[epoch]=0.5*np.mean((output-y)**2)
            errors = output - y.T
            self.weights -= ((self.learning_rate * (X.T.dot(errors)))/np.size
            self.bias -= (self.learning_rate * np.mean(output))
            if (error[epoch]==0):break
        return self.weights ,self.bias

    def predict(self, X):
        net_input = np.dot(X, self.weights) + self.bias
        output = self.activation(net_input)
        return np.where(output >= 0.5, 1, 0)

    def activation(self, X):
        return X
```

```
In [4]: def accuracy(y_true, y_pred):
    accuracy = np.sum(y_true == y_pred) / len(y_true)
    return accuracy
```

```
In [5]: def MSE(y_pred, y):
    return np.mean((y_pred - y)**2)
```

```
In [6]: df = pd.read_csv("data_banknote_authentication.csv")
df
```

Out[6]:

	variance	skewness	curtosis	entropy	class
0	3.62160	8.66610	-2.8073	-0.44699	0
1	4.54590	8.16740	-2.4586	-1.46210	0
2	3.86600	-2.63830	1.9242	0.10645	0
3	3.45660	9.52280	-4.0112	-3.59440	0
4	0.32924	-4.45520	4.5718	-0.98880	0
...
1367	0.40614	1.34920	-1.4501	-0.55949	1
1368	-1.38870	-4.87730	6.4774	0.34179	1
1369	-3.75030	-13.45860	17.5932	-2.77710	1
1370	-3.56370	-8.38270	12.3930	-1.28230	1
1371	-2.54190	-0.65804	2.6842	1.19520	1

1372 rows × 5 columns

In [7]: *# separate the independent and dependent variables*

```
X_data = df.iloc[ : , 0 : 2]
target = df.iloc[ : , -1]
print(X_data)
print("-"*50)
print(target)
```

```
      variance  skewness
0      3.62160   8.66610
1      4.54590   8.16740
2      3.86600  -2.63830
3      3.45660   9.52280
4      0.32924  -4.45520
...
1367    0.40614   1.34920
1368   -1.38870  -4.87730
1369   -3.75030 -13.45860
1370   -3.56370  -8.38270
1371   -2.54190  -0.65804
```

[1372 rows x 2 columns]

```
-----
0      0
1      0
2      0
3      0
4      0
..
1367    1
1368    1
1369    1
1370    1
1371    1
```

Name: class, Length: 1372, dtype: int64

In [8]: `scale= StandardScaler()`
standardization of independent variables
`scaled_data = scale.fit_transform(X_data)`
`print(scaled_data)`

```
[[ 1.12180565  1.14945512]
 [ 1.44706568  1.06445293]
 [ 1.20780971 -0.77735215]
 ...
 [-1.47235682 -2.62164576]
 [-1.40669251 -1.75647104]
 [-1.04712236 -0.43982168]]
```

```
In [9]: print(target)
```

```
0      0
1      0
2      0
3      0
4      0
..
1367    1
1368    1
1369    1
1370    1
1371    1
Name: class, Length: 1372, dtype: int64
```


```
In [10]: scaled_data = pd.DataFrame(scaled_data)
print(scaled_data)
```

```
      0      1
0  1.121806  1.149455
1  1.447066  1.064453
2  1.207810 -0.777352
3  1.063742  1.295478
4 -0.036772 -1.087038
...
1367 -0.009711 -0.097693
1368 -0.641313 -1.158984
1369 -1.472357 -2.621646
1370 -1.406693 -1.756471
1371 -1.047122 -0.439822

[1372 rows x 2 columns]
```



```
In [11]: plt.scatter(scaled_data.iloc[:, 0], scaled_data.iloc[:, 1], c=target)
plt.show()
```



```
In [12]: X_train, X_test, y_train, y_test = train_test_split(scaled_data, target , test_size=0.2)
```

Perceptron


```
In [13]: percept = Perceptron(learning_rate=0.0001, epochs=1000)
percept.fit(X_train, y_train)
percept_predictions = percept.predict(X_test)
```

```
In [14]: print("Perceptron classification accuracy:", accuracy(y_test, percept_predictions))
print("-"*50)
print("Perceptron classification MSE:", MSE(percept_predictions , y_test)*100)
```

Perceptron classification accuracy: 86.9090909090909

Perceptron classification MSE: 13.090909090909092

```
In [15]: X_train=X_train.to_numpy()
fig = plt.figure()
ax = fig.add_subplot(1, 1, 1)
plt.scatter(X_train[:, 0], X_train[:, 1], marker="o", c=y_train)
x0_1 = np.amin(X_train[:, 0])
x0_2 = np.amax(X_train[:, 0])
x1_1 = (-percept.weights[0] * x0_1 - percept.bias) / percept.weights[1]
x1_2 = (-percept.weights[0] * x0_2 - percept.bias) / percept.weights[1]
ax.plot([x0_1, x0_2], [x1_1, x1_2], "k")
ymin = np.amin(X_train[:, 1])
ymax = np.amax(X_train[:, 1])
ax.set_ylim([ymin - 3, ymax + 3])
plt.ylim(-3,3)
plt.show()
```



Adaline


```
In [16]: ada = Adaline()
w, b = ada.fit(X_train,y_train)
ada_predictions = ada.predict(X_test)
```

```
In [17]: print("Adaline classification accuracy:", accuracy(y_test, ada_predictions)*100)
print("-"*50)
print("Adaline classification MSE:", MSE(ada_predictions , y_test)*100)
```

Adaline classification accuracy: 69.45454545454545

Adaline classification MSE: 30.545454545454547

```
In [18]: fig = plt.figure()
ax = fig.add_subplot(1, 1, 1)
plt.scatter(X_train[:, 0], X_train[:, 1], marker="o", c=y_train)
x0_1 = np.amin(X_train[:, 0])
x0_2 = np.amax(X_train[:, 0])
x1_1 = (-w[0] * x0_1 - b) / w[1]
x1_2 = (-w[0] * x0_2 - b) / w[1]
ax.plot([x0_1, x0_2], [x1_1, x1_2], "k")
ymin = np.amin(X_train[:, 1])
ymax = np.amax(X_train[:, 1])
ax.set_ylim([ymin - 3, ymax + 3])
plt.ylim(-3,3)
plt.show()
```



In []:

In []:

In []: