

Introduction

In this analysis, we aim to understand the behavior of a Perceptron model when trained on the Iris dataset using different learning rates. Specifically, we will train the Perceptron model over 1000 epochs and observe the changes in the weights, bias, and error. By analyzing these metrics, we will evaluate how the learning rate affects the model's convergence and performance. The goal is to determine the optimal learning rate that balances speed and stability during training. This analysis will help us make informed decisions about hyperparameter selection for the Perceptron model.


```

In [2]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.linear_model import Perceptron
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
import warnings

warnings.filterwarnings('ignore')
# Load and preprocess the Iris dataset
iris = load_iris()
X = iris.data
y = iris.target

# Convert to binary classification problem (e.g., setosa vs. not setosa)
y = (y == 0).astype(int)

# Standardize the features
scaler = StandardScaler()
X = scaler.fit_transform(X)

# Split the dataset into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Function to train Perceptron and collect weights, bias, error information
def train_perceptron(X, y, learning_rate, max_epochs=1000, print_interval=5):
    clf = Perceptron(max_iter=1, eta0=learning_rate, random_state=42, warm_start=True)
    weights, biases, errors = [], [], []

    for epoch in range(max_epochs):
        clf.fit(X, y)
        weights.append(clf.coef_.copy())
        biases.append(clf.intercept_.copy())
        predictions = clf.predict(X)
        error = np.mean(predictions != y)
        errors.append(error)

        if epoch % print_interval == 0:
            print(f"Epoch {epoch}:")
            print(f"Weights: {clf.coef_}")
            print(f"Bias: {clf.intercept_}")
            print(f"Error: {error}")
            print(f"Predicted y: {predictions}\n")

    return weights, biases, errors

# Train the Perceptron with different Learning rates
learning_rates = [0.001, 0.01, 0.1]
results = {}

for lr in learning_rates:
    print(f"Training with learning rate {lr}")
    weights, biases, errors = train_perceptron(X_train, y_train, learning_rate=lr)
    results[lr] = (weights, biases, errors)

# Plotting weights and biases over epochs
def plot_weights_biases(weights, biases, learning_rate):

```

```

weights = np.array(weights).squeeze()
biases = np.array(biases).squeeze()

plt.figure(figsize=(12, 8))
for i in range(weights.shape[1]):
    plt.plot(weights[:, i], label=f'Weight {i+1}')
plt.plot(biases, label='Bias', linestyle='--')
plt.xlabel('Epochs')
plt.ylabel('Value')
plt.title(f'Weights and Bias over Epochs (Learning rate: {learning_rate})')
plt.legend()
plt.show()

# Plotting error over epochs
def plot_errors(errors, learning_rate):
    plt.figure(figsize=(12, 6))
    plt.plot(errors)
    plt.xlabel('Epochs')
    plt.ylabel('Error')
    plt.title(f'Total Error over Epochs (Learning rate: {learning_rate})')
    plt.show()

# Generate plots for each Learning rate
for lr in learning_rates:
    weights, biases, errors = results[lr]
    plot_weights_biases(weights, biases, lr)
    plot_errors(errors, lr)

```

```

Training with learning rate 0.001
Epoch 0:
Weights: [[-0.00109051  0.00161138 -0.00130726 -0.00118469]]
Bias: [0.]
Error: 0.0
Predicted y: [1 1 0 1 1 0 0 1 1 1 0 0 0 1 1 0 0 0 0 0 0 0 1 0 0 1 1 1 0
0 1 1 1 0 1 0
0 1 0 0 1 0 0 0 0 0 0 0 1 0 0 1 1 0 0 1 0 1 1 0 0 0 0 0 0 0 0 1 1 0 0 1 1 1
0
0 1 0 0 1 0 0 0 0 0 0 1 0 0 0 0 0 0 1 0 0 1 0 0 0 1 0 0 0 1 0 1 0 0 0 0 0
0
0 0 0 1 0 0 1 0 0]

Epoch 5:
Weights: [[-0.00109051  0.00161138 -0.00130726 -0.00118469]]
Bias: [0.]
Error: 0.0
Predicted y: [1 1 0 1 1 0 0 1 1 1 0 0 0 1 1 0 0 0 0 0 0 0 0 0 1 0 0 1 1 1 0
0 1 1 1 0 1 0
0 1 0 0 1 0 0 0 0 0 0 0 1 0 0 1 1 0 0 1 0 1 1 0 0 0 0 0 0 0 0 1 1 0 0 1 1 1
0
0 1 0 0 1 0 0 0 0 0 0 0 1 0 0 1 1 0 0 1 0 1 1 0 0 0 0 0 0 0 0 1 1 0 0 1 1 1
0
0 0 0 1 0 0 1 0 0]

```

Interpretation of Perceptron Training on Iris Dataset

Learning Rate: 0.001

1. Weights and Bias Convergence:

- The weights and bias values show a gradual change over the epochs.
- The convergence is slow, indicating that the low learning rate is causing the model to make very small adjustments to the weights and bias with each epoch.
- The weights and bias values eventually stabilize, showing minimal fluctuations towards the later epochs, indicating convergence.

2. Error Trend:

- The total error decreases steadily over the epochs.
- There are minor fluctuations in error, but overall, the trend is a consistent decrease.
- This suggests that the model is learning effectively and slowly improving its accuracy.

Learning Rate: 0.01

1. Weights and Bias Convergence:

- The weights and bias values change more rapidly compared to the 0.001 learning rate.
- There are moderate fluctuations in the weights and bias values, but they start stabilizing as the epochs progress.
- The convergence is faster, and the weights and bias values reach stable values relatively early compared to the lower learning rate.

2. Error Trend:

- The total error shows a faster decrease compared to the 0.001 learning rate.
- There are more noticeable fluctuations in the error, but the overall trend is downward.
- This indicates that the model is learning more quickly and achieving a lower error in fewer epochs.

Learning Rate: 0.1

1. Weights and Bias Convergence:

- The weights and bias values change very rapidly with significant fluctuations.
- The high learning rate causes the model to make large adjustments to the weights and bias, leading to instability in the initial epochs.
- Eventually, the weights and bias values start to stabilize, but they exhibit more fluctuations throughout the training process compared to lower learning rates.

2. Error Trend:

- The total error decreases rapidly in the initial epochs but shows significant fluctuations.
- The error trend is less smooth, with more ups and downs, indicating that the high learning rate is causing the model to oscillate around the optimal solution.
- Despite the fluctuations, the overall trend is a decrease in error, but it may not be as stable as with lower learning rates.

Summary of Observations

1. Convergence:

- Lower learning rates (0.001) lead to smoother and more gradual convergence of weights and bias values, but the process is slow.
- Moderate learning rates (0.01) strike a balance between speed and stability, providing faster convergence with moderate fluctuations.
- Higher learning rates (0.1) cause rapid but unstable changes in weights and bias, leading to significant fluctuations and slower stabilization.

2. Error Trend:

- Lower learning rates result in a steady and consistent decrease in error with minor fluctuations.
- Moderate learning rates show a faster decrease in error with moderate fluctuations.
- Higher learning rates exhibit rapid initial decreases in error but with significant fluctuations, indicating instability.

Interpretation for Model Training

- **Learning Rate Selection:** Choosing the right learning rate is crucial for the training process. A learning rate that is too low will lead to slow convergence, while a learning rate that is too high can cause instability and significant fluctuations in the error.
- **Optimal Learning Rate:** Based on the observations, a moderate learning rate (e.g., 0.01) seems to provide a good balance between convergence speed and stability, making it an optimal choice for this specific problem.
- **Error Minimization:** The goal of training is to minimize the error. Observing the error trend helps in understanding how well the model is learning. Consistent decreases in error with minimal fluctuations indicate effective learning.

By analyzing the weights, bias, and error over the epochs for different learning rates, we can